



CSYE 7105

Team 13

Parallelizing Plant Health

**A Comparative Study of Parallel Deep Learning
for Plant Disease Classification**

**He Zhou
Weiyu Wang**

Outline

Introduction

1

- ✓ Background
- ✓ Motivation
- ✓ Goals

Methodology

2

- ✓ Roadmap
- ✓ Mechanism
- ✓ Tools

Dataset Description

3

- ✓ Overview
- ✓ Splitting
- ✓ EDA

Result & Analysis

4

- ✓ Pre-processing
- ✓ Pre-training
- ✓ Hyperparameter tuning
- ✓ Final Evaluation

Conclusion

5

- ✓ How does parallelism perform in each phase?

1 Introduction

1.1 Background

How much economic loss
do plant diseases cause
each year?

1 Introduction

1.1 Background

\$220 billion annually

according to the Food and Agriculture Organization of the **United Nations**

developing countries, small stakeholders...

1 Introduction

1.2 Motivation & goal

Motivation

- **Achieve timely disease identification**
- **Datasets are huge, Trainings are slow** - facing the challenge of the large image datasets

Goal

- **Accelerate** the training by leveraging the power of parallel techniques
- **Compare** the time efficiency of serial and parallel deep learning, analyze the potential benefits of parallel approaches.

2 Methodology

Roadmap

Step	Roadmap	Focuses	Tools	Parallelization Comparison
1	Exploratory Data Analysis	Number of classes	NumPy, Matplotlib	
2	Data Preprocessing	Mean, std calculation for dataset normalization	(1) Dask Scheduler (2) PyTorch	√
3	Model Definition	Resnet18	Torchvision	
4	Model Pretraining	(1) Train on an initial set of hyperparameters on small number of epochs (2) Compare time efficiency between serial and parallel techniques	PyTorch (1) Serial (2) Using multi-process (num_workers) (3) Using multi-thread (DataParallel)	√
5	Hyperparameter Tuning	Bayesian optimization	Optuna	√
6	Final Training and Model Evaluation	Use optimized hyperparameters on extended epochs accuracy, precision, recall, and F1 score, confusion matrix	PyTorch Scikit-learn, Seaborn	

3 Dataset Description

3.1 Overview



SAMIR BHATTARAI · UPDATED 5 YEARS AGO



860

New Notebook



Download (3 GB)



New Plant Diseases Dataset

Image dataset containing different healthy and unhealthy crop leaves.



Data Card

Code (280)

Discussion (4)

About Dataset

This dataset is recreated using offline augmentation from the original dataset. The original dataset can be found on [this github repo](#). This dataset consists of about 87K rgb images of healthy and diseased crop leaves which is categorized into 38 different classes. The total dataset is divided into 80/20 ratio of training and validation set preserving the directory structure. A new directory containing 33 test images is created later for prediction purpose.

Usability ⓘ

7.50

License

Data files © Original Authors

Expected update frequency

Not specified

Tags

Biology

Image

Source:

<https://www.kaggle.com/datasets/vipooooool/new-plant-diseases-dataset>

Total size:


1.43 GB, contains 87.9k images

Dataset after splitting:

Subset	Image counts	Percent
train	61494	70%
valid	17572	20%
test	8814	10%
Total	87880	100%

3 Dataset Description

3.2 Exploratory Data Analysis



3 Dataset Description

3.2 Exploratory Data Analysis

The dataset is organized into 38 distinct categories based on plant species and specific diseases. Among these categories, there are 14 distinct plant categories and 26 types of plant diseases.

14 Plant Categories: Apple, Blueberry, Cherry, Corn, Grape, Orange, Peach, Pepper, Potato, Raspberry, Soybean, Squash, Strawberry, Tomato.

	File Name	Image Count
0	Apple__Apple_scab	1764
1	Apple__Black_rot	1738
2	Apple__Cedar_apple_rust	1540
3	Apple__healthy	1757
4	Blueberry__healthy	1589
5	Cherry_(including_sour)__Powdery_mildew	1472
6	Cherry_(including_sour)__healthy	1597
7	Corn_(maize)__Cercospora_leaf_spot Gray_leaf_...	1436
8	Corn_(maize)__Common_rust_	1668
9	Corn_(maize)__Northern_Leaf_Blight	1669
10	Corn_(maize)__healthy	1626
11	Grape__Black_rot	1652
12	Grape__Esca_(Black_Measles)	1680
13	Grape__Leaf_blight_(Isariopsis_Leaf_Spot)	1506
14	Grape__healthy	1480
15	Orange__Haunglongbing_(Citrus_greening)	1758
16	Peach__Bacterial_spot	1608
17	Peach__healthy	1512
18	Pepper,_bell__Bacterial_spot	1673

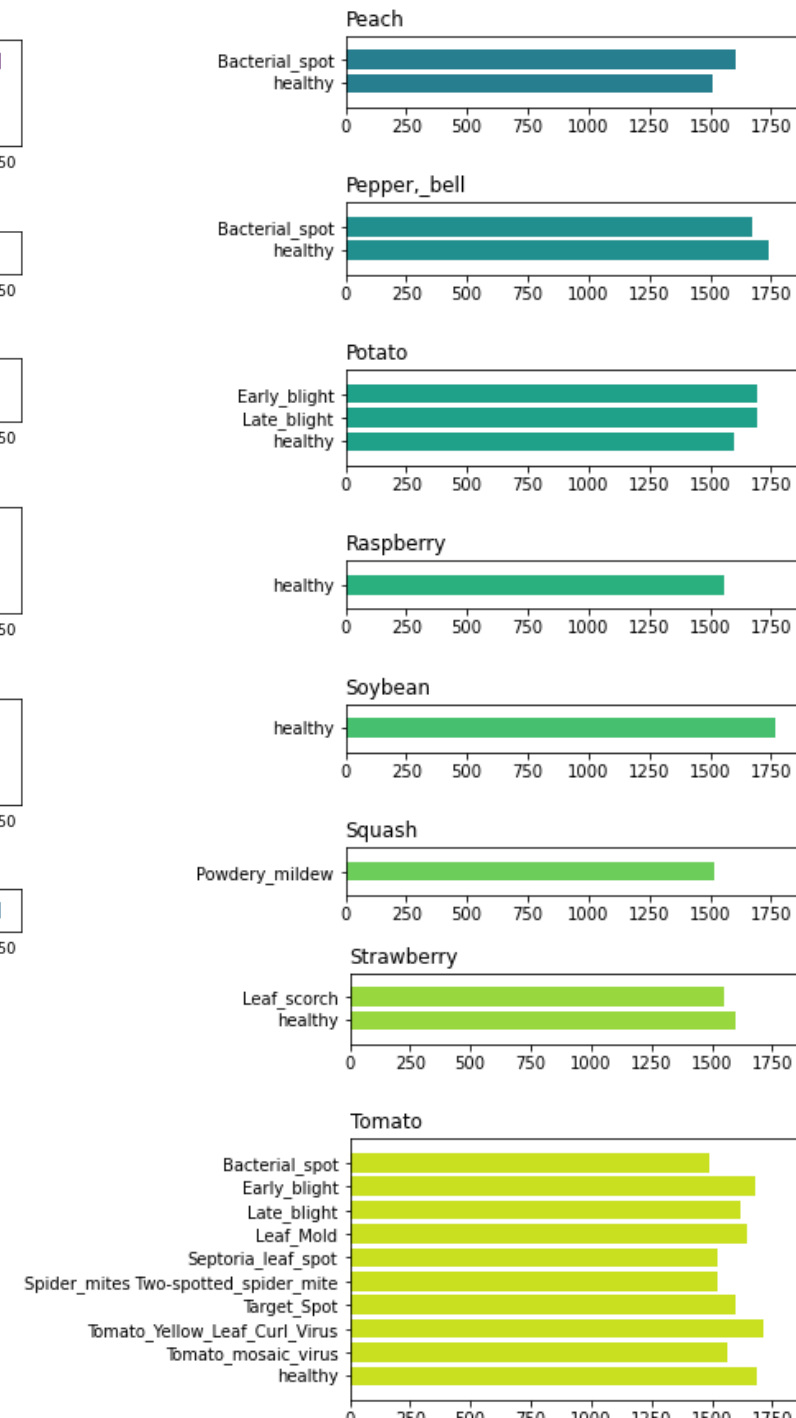
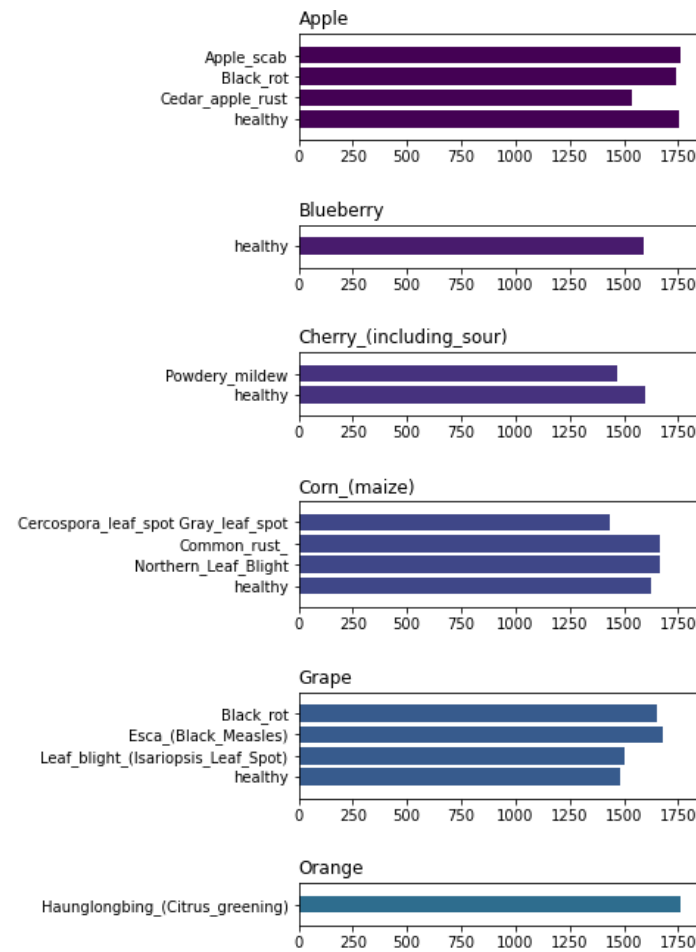
19	Pepper,_bell__healthy	1739
20	Potato__Early_blight	1696
21	Potato__Late_blight	1696
22	Potato__healthy	1596
23	Raspberry__healthy	1558
24	Soybean__healthy	1769
25	Squash__Powdery_mildew	1519
26	Strawberry__Leaf_scorch	1552
27	Strawberry__healthy	1596
28	Tomato__Bacterial_spot	1489
29	Tomato__Early_blight	1680
30	Tomato__Late_blight	1619
31	Tomato__Leaf_Mold	1646
32	Tomato__Septoria_leaf_spot	1526
33	Tomato__Spider_mites Two-spotted_spider_mite	1523
34	Tomato__Target_Spot	1598
35	Tomato__Tomato_Yellow_Leaf_Curl_Virus	1715
36	Tomato__Tomato_mosaic_virus	1566
37	Tomato__healthy	1685

3 Dataset Description

3.2 Exploratory Data Analysis

The dataset is organized into 38 distinct categories based on plant species and specific diseases. Among these categories, there are 14 distinct plant categories and 26 types of plant diseases.

14 Plant Categories: Apple, Blueberry, Cherry, Corn, Grape, Orange, Peach, Pepper, Potato, Raspberry, Soybean, Squash, Strawberry, Tomato.



4 Result & Analysis

4.1 Data Preprocessing

Data Cleaning: Identify or delete dirty image data and control image data pixel consistency.

Import all the crucial libraries:

```
import os
import dask
import dask.array as da from skimage.io
import imread import numpy as np
```

Data Transformation:

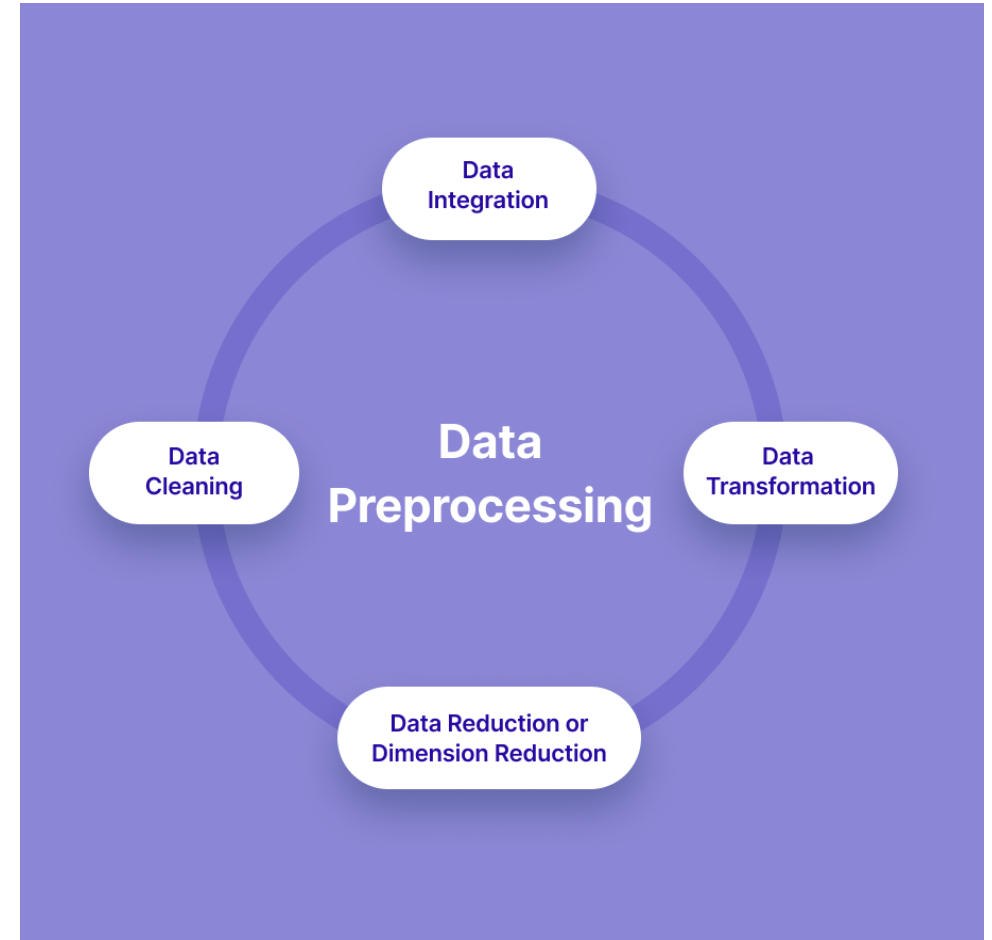
Read the image and convert it to an array of values.
Convert an image to a PyTorch tensor using transforms.Compose.

Data Normalization: Calculate mean and Standard Deviation

Classic ImageNet statistics:

mean = [0.485, 0.456, 0.406]

std = [0.229, 0.224, 0.225]

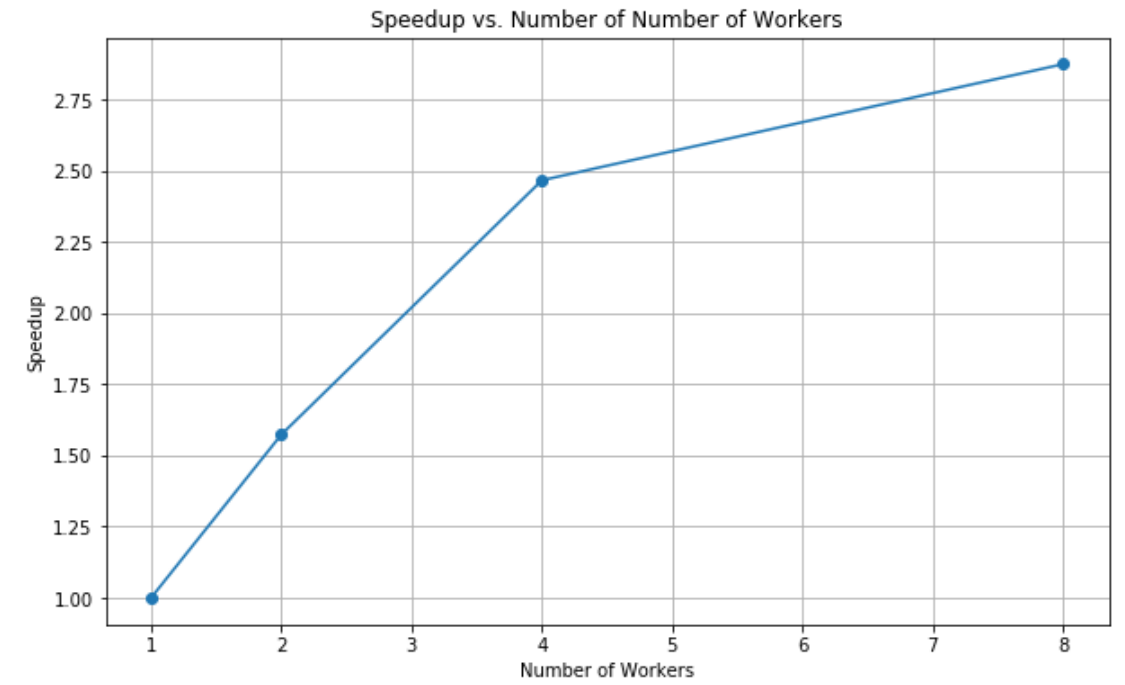



4 Result & Analysis

4.1 Data Preprocessing

Using Dask

Workers	Time(s)	Speed-up
1	1321.3	1
2	840.0	1.57
4	526.0	2.47
8	459.9	2.87





4 Result & Analysis

4.1 Data Preprocessing

Using PyTorch

Workers	Time(s)	Speed-up
1	378.3	1
2	207.8	1.82
4	122.3	3.09
8	83.7	4.52



4 Result & Analysis

4.2 Pre-training

Serial Training

Hardware

CPU

Intel(R) Xeon(R) Gold 6240 CPU @
2.60GHz
Architecture: X86_64
Count: 4 (1 processes)

GPU

Tesla T4
Count: 1

NVIDIA-SMI 535.104.05				Driver Version: 535.104.05		CUDA Version: 12.2	
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.
							MIG M.
0	Tesla T4		Off	00000000:3B:00.0	Off		0
N/A	52C	P0	29W / 70W	2MiB / 15360MiB		6%	Default
							N/A
Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
	ID	ID				Usage	
No running processes found							

4 Result & Analysis

4.2 Pre-training

Serial Training

Progress

```
Running on 1 GPU with 1 process-----
Epoch [1/5]-----
| Train Loss: 2.6422 | Train Acc: 0.3406 | Val Loss: 2.2021, Val Acc: 0.4074
| Elapsed Time: 445.8777 s | Max GPU Memory Alloc: 1960.8457 MB
Epoch [2/5]-----
| Train Loss: 2.0209 | Train Acc: 0.4719 | Val Loss: 1.7971, Val Acc: 0.5216
| Elapsed Time: 424.5686 s | Max GPU Memory Alloc: 1960.8472 MB
Epoch [3/5]-----
| Train Loss: 1.8274 | Train Acc: 0.5141 | Val Loss: 1.6481, Val Acc: 0.5436
| Elapsed Time: 425.9290 s | Max GPU Memory Alloc: 1960.8481 MB
Epoch [4/5]-----
| Train Loss: 1.2209 | Train Acc: 0.6451 | Val Loss: 1.2316, Val Acc: 0.6412
| Elapsed Time: 428.5931 s | Max GPU Memory Alloc: 1960.8491 MB
Epoch [5/5]-----
| Train Loss: 1.2085 | Train Acc: 0.6483 | Val Loss: 1.2196, Val Acc: 0.6451
| Elapsed Time: 427.7931 s | Max GPU Memory Alloc: 1960.8501 MB

total_time: 2153.4934573173523
best_valid_acc: 0.6450603232415206
```

4 Result & Analysis

4.2 Pre-training

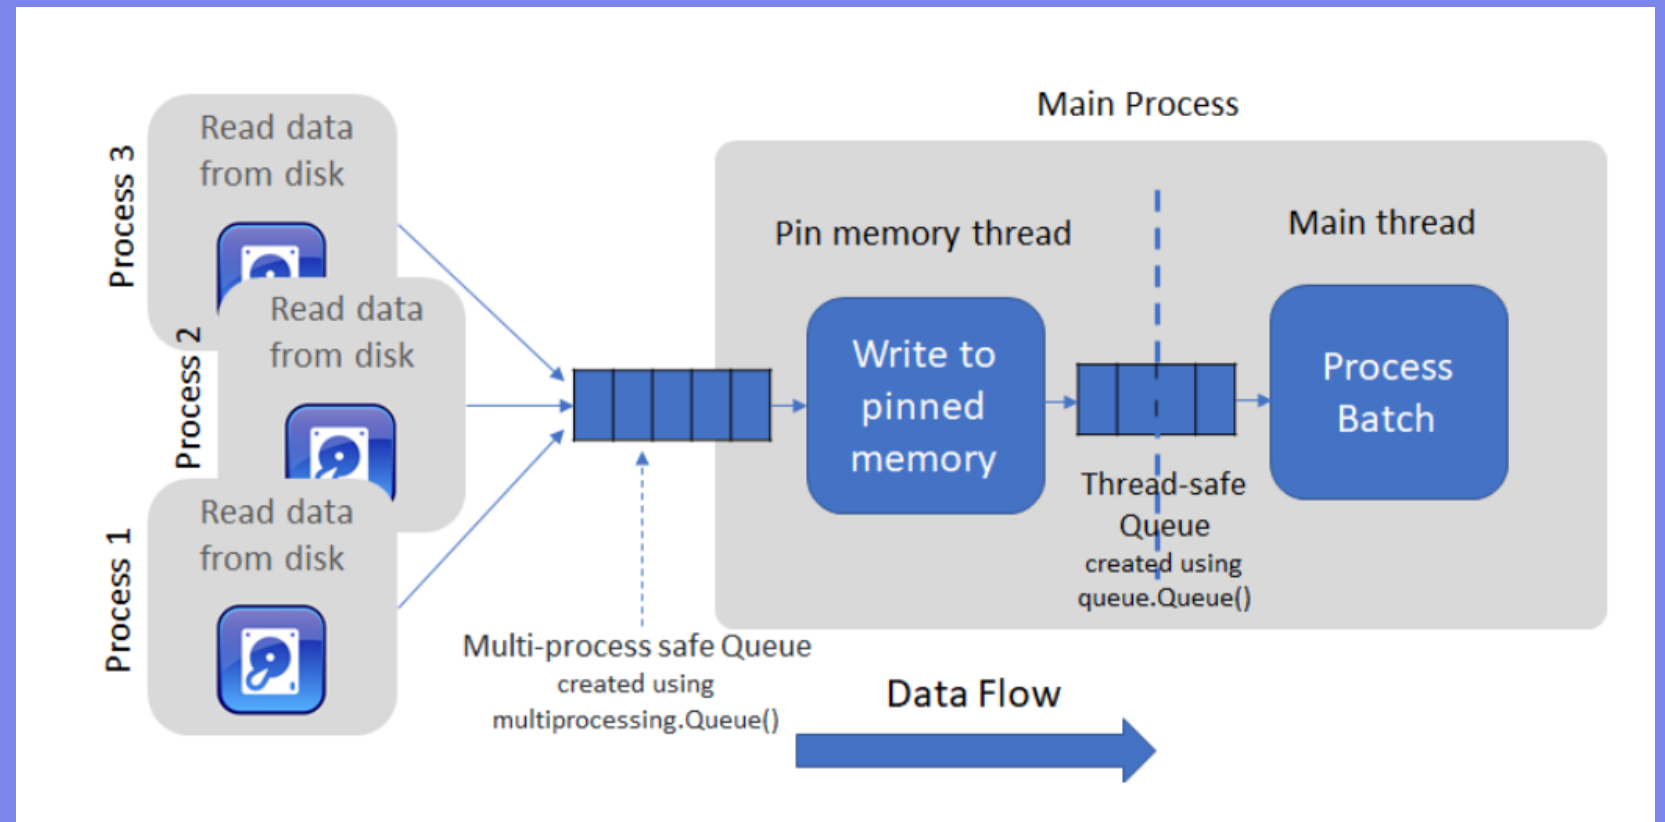
Using Multi-process (workers)

Parallel Mechanism

Multi workers in DataLoader

It uses multiple processes to load and process batch data from the disk into memory and puts it into a queue.

The main training process then consumes batches from this queue, reducing the overall time taken for each training iteration.



4 Result & Analysis

4.2 Pre-training

Using Multi-process (workers)

Hardware

CPU

Intel(R) Xeon(R) Gold 6240 CPU @
2.60GHz
Architecture: X86_64
Count: 4 (2, 4, 8 processes)

GPU

Tesla T4
Count: 1

NVIDIA-SMI 535.104.05				Driver Version: 535.104.05		CUDA Version: 12.2	
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.
							MIG M.
0	Tesla T4		Off	00000000:3B:00.0	Off		0
N/A	52C	P0	29W / 70W	2MiB / 15360MiB		6%	Default
							N/A
Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
	ID	ID				Usage	
No running processes found							

4 Result & Analysis

4.2 Pre-training

Using Multi-process (workers)

2

processes

Progress

```
Running on 1 GPU with 2 processes-----
Epoch [1/5]-----
| Train Loss: 2.6422 | Train Acc: 0.3406 | Val Loss: 2.2021, Val Acc: 0.4074
| Elapsed Time: 248.3875 s | Max GPU Memory Alloc: 1960.8457 MB
Epoch [2/5]-----
| Train Loss: 2.0209 | Train Acc: 0.4719 | Val Loss: 1.7971, Val Acc: 0.5216
| Elapsed Time: 227.4754 s | Max GPU Memory Alloc: 1960.8472 MB
Epoch [3/5]-----
| Train Loss: 1.8274 | Train Acc: 0.5141 | Val Loss: 1.6481, Val Acc: 0.5436
| Elapsed Time: 231.8289 s | Max GPU Memory Alloc: 1960.8481 MB
Epoch [4/5]-----
| Train Loss: 1.2209 | Train Acc: 0.6451 | Val Loss: 1.2316, Val Acc: 0.6412
| Elapsed Time: 227.1396 s | Max GPU Memory Alloc: 1960.8491 MB
Epoch [5/5]-----
| Train Loss: 1.2085 | Train Acc: 0.6483 | Val Loss: 1.2196, Val Acc: 0.6451
| Elapsed Time: 234.1333 s | Max GPU Memory Alloc: 1960.8501 MB

total_time: 1169.8230679035187
best_valid_acc: 0.6450603232415206
```


4 Result & Analysis

4.2 Pre-training

Using Multi-process (workers)

4

processes

Progress

```
Running on 1 GPU with 4 processes-----  
Epoch [1/5]-----  
| Train Loss: 2.6422 | Train Acc: 0.3406 | Val Loss: 2.2021, Val Acc: 0.4074  
| Elapsed Time: 223.4081 s | Max GPU Memory Alloc: 1960.8457 MB  
Epoch [2/5]-----  
| Train Loss: 2.0209 | Train Acc: 0.4719 | Val Loss: 1.7971, Val Acc: 0.5216  
| Elapsed Time: 222.7155 s | Max GPU Memory Alloc: 1960.8472 MB  
Epoch [3/5]-----  
| Train Loss: 1.8274 | Train Acc: 0.5141 | Val Loss: 1.6481, Val Acc: 0.5436  
| Elapsed Time: 223.5448 s | Max GPU Memory Alloc: 1960.8481 MB  
Epoch [4/5]-----  
| Train Loss: 1.2209 | Train Acc: 0.6451 | Val Loss: 1.2316, Val Acc: 0.6412  
| Elapsed Time: 220.5944 s | Max GPU Memory Alloc: 1960.8491 MB  
Epoch [5/5]-----  
| Train Loss: 1.2085 | Train Acc: 0.6483 | Val Loss: 1.2196, Val Acc: 0.6451  
| Elapsed Time: 223.2077 s | Max GPU Memory Alloc: 1960.8501 MB  
  
total_time: 1114.2051734924316  
best_valid_acc: 0.6450603232415206
```

4 Result & Analysis

4.2 Pre-training

Using Multi-process (workers)

Progress

```
Running on 1 GPU with 8 processes-----
Epoch [1/5]-----
| Train Loss: 2.6422 | Train Acc: 0.3406 | Val Loss: 2.2021, Val Acc: 0.4074
| Elapsed Time: 224.6506 s | Max GPU Memory Alloc: 1960.8457 MB
Epoch [2/5]-----
| Train Loss: 2.0209 | Train Acc: 0.4719 | Val Loss: 1.7971, Val Acc: 0.5216
| Elapsed Time: 222.2073 s | Max GPU Memory Alloc: 1960.8472 MB
Epoch [3/5]-----
| Train Loss: 1.8274 | Train Acc: 0.5141 | Val Loss: 1.6481, Val Acc: 0.5436
| Elapsed Time: 222.1483 s | Max GPU Memory Alloc: 1960.8481 MB
Epoch [4/5]-----
| Train Loss: 1.2209 | Train Acc: 0.6451 | Val Loss: 1.2316, Val Acc: 0.6412
| Elapsed Time: 222.4511 s | Max GPU Memory Alloc: 1960.8491 MB
Epoch [5/5]-----
| Train Loss: 1.2085 | Train Acc: 0.6483 | Val Loss: 1.2196, Val Acc: 0.6451
| Elapsed Time: 221.2881 s | Max GPU Memory Alloc: 1960.8501 MB

total_time: 1113.5023527145386
best_valid_acc: 0.6450603232415206
```

Warning

local/lib/python3.9/site-packages/torch/utils/data/dataloader.py:557: UserWarning: This DataLoader will create 8 worker processes in total. Our suggested max number of worker in current system is 4, which is smaller than what this DataLoader is going to create. Please be aware that excessive worker creation might get DataLoader running slow or even freeze, lower the worker number to avoid potential slowness/freeze if necessary.

8

processes

over-subscription

4 Result & Analysis


4.2 Pre-training

Using Multi-process

Comparison

Count of Process	Elapsed time (s)	Speedup
Serial (1 process)	2153.4935	1
2 processes	1169.8231	1.8409
4 processes	1114.2052	1.9328
8 processes	1113.5024	1.9340

Multi-process Elapsed Time and Speedup

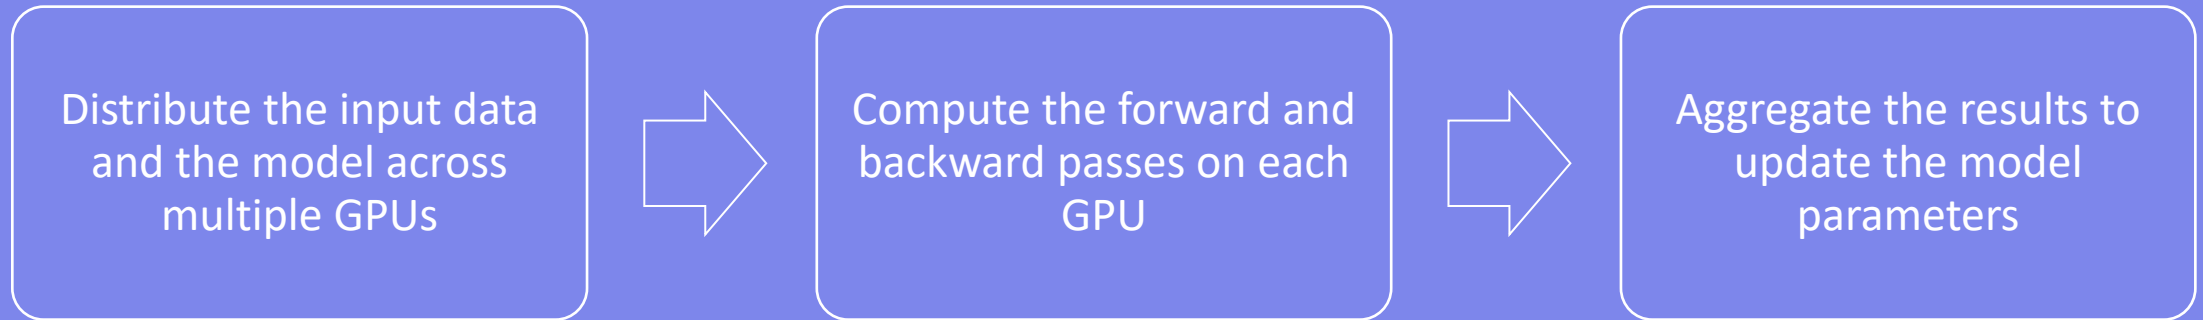


4 Result & Analysis

4.2 Pre-training

Using Multi-thread (DataParallel)

Parallel Mechanism



Before we dive in, let's clarify why, despite the added complexity, you would consider using `DistributedDataParallel` over `DataParallel`:

- First, `DataParallel` is single-process, multi-thread, and only works on a single machine, while

`DistributedDataParallel` is multi-process and works for both single- and multi- machine training. `DataParallel` is usually slower than `DistributedDataParallel` even on a single machine due to GIL contention across threads, per-

4 Result & Analysis

4.2 Pre-training

Using Multi-thread (DataParallel)

Hardware

CPU

Intel(R) Xeon(R) Gold 6240 CPU @ 2.60GHz

Architecture: X86_64

Count: 4

GPU

Tesla T4

Count: 2, 4(expected)

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
NVIDIA-SMI 470.161.03 Driver Version: 470.161.03 CUDA Version: 11.4													
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
GPU		Name		Persistence-M		Bus-Id		Disp.A		Volatile Uncorr. ECC			
Fan		Temp		Perf		Pwr:Usage/Cap		Memory-Usage		GPU-Util Compute M.			
										MIG M.			
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
0		Tesla T4		Off		00000000:00:04.0		Off		0			
N/A		46C		P8		10W / 70W		0MiB / 15109MiB		0%			
										Default			
										N/A			
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
1		Tesla T4		Off		00000000:00:05.0		Off		0			
N/A		43C		P8		10W / 70W		0MiB / 15109MiB		0%			
										Default			
										N/A			
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
Processes:													
GPU		GI		CI		PID		Type		Process name		GPU Memory	
		ID		ID								Usage	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													
No running processes found													
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+													

4 Result & Analysis

4.2 Pre-training

Using Multi-thread (DataParallel)

GPU
2

Progress

```
Running DataParallel on 2 GPUs-----
Epoch [1/5]-----
| Train Loss: 2.7000 | Train Acc: 0.3273 | Val Loss: 2.3335, Val Acc: 0.3904
| Elapsed Time: 335.0122 s | Max GPU Memory Alloc: 1079.4219 MB
Epoch [2/5]-----
| Train Loss: 2.0822 | Train Acc: 0.4565 | Val Loss: 1.8347, Val Acc: 0.5254
| Elapsed Time: 321.0952 s | Max GPU Memory Alloc: 1079.4233 MB
Epoch [3/5]-----
| Train Loss: 1.9114 | Train Acc: 0.4969 | Val Loss: 1.7686, Val Acc: 0.5411
| Elapsed Time: 319.5789 s | Max GPU Memory Alloc: 1079.4243 MB
Epoch [4/5]-----
| Train Loss: 1.2854 | Train Acc: 0.6260 | Val Loss: 1.2389, Val Acc: 0.6403
| Elapsed Time: 324.5508 s | Max GPU Memory Alloc: 1079.4253 MB
Epoch [5/5]-----
| Train Loss: 1.2701 | Train Acc: 0.6299 | Val Loss: 1.2242, Val Acc: 0.6469
| Elapsed Time: 329.6849 s | Max GPU Memory Alloc: 1079.4263 MB

total_time: 1630.4430594444275
best_valid_acc: 0.6469383109492375
```


4 Result & Analysis

4.2 Pre-training

Using Multi-thread (DataParallel)

Comparison

Count of GPU	Elapsed time (s)	Speedup
Serial (1 GPU)	2153.4935	1
2 GPU	1630.4431	1.3208
4 GPU	-	-



4 Result & Analysis

4.2 Pre-training

Comparison Summary



Training	Resources	Elapsed Time (s)	Speedup
Serial	Main process	2153.4935	1
Multi-process (workers)	2 processes	1169.8231	1.8409
	4 processes	1114.2052	1.9328
	8 processes	1113.5024	1.9340
Multi-thread (DataParallel)	2 GPUs	1630.4431	1.3208
	4 GPUs	-	-

DistributedDataParallel

4 Result & Analysis

4.2 Pre-training

Comparison Summary



4 Result & Analysis

4.3 Hyperparameter Tuning

Parallel Mechanism

Optuna study
Apply parallelism using
Joblib threading

Hyperparameters we tuned:

- learning rate
- step size
- weight decay
- momentum

```
import joblib
from joblib import delayed
from joblib import Parallel
```

```
with Parallel(n_jobs=n_jobs, prefer="threads") as parallel:
    if not isinstance(
        parallel._backend, joblib.parallel.ThreadingBackend
    ) and isinstance(self._storage, storages.InMemoryStorage):
        msg = (
            "The default storage cannot be shared by multiple processes. "
            "Please use an RDB (RDBStorage) when you use joblib for "
            "multi-processing. The usage of RDBStorage can be found in "
            "https://optuna.readthedocs.io/en/stable/tutorial/rdb.html."
```


4.3 Hyperparameter Tuning

CPU

GPU

A100-SXM4-80GB
Count: 1

NVIDIA-SMI 535.104.05			Driver Version: 535.104.05		CUDA Version: 12.2	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.
						MIG M.
0	NVIDIA A100-SXM4-80GB	Off	00000000:81:00.0	Off		0
N/A	39C	P0	59W / 500W	4MiB / 81920MiB	0%	Default Disabled
Processes:						
GPU	GI	CI	PID	Type	Process name	GPU Memory Usage
	ID	ID				
No running processes found						

4 Result & Analysis

4.3 Hyperparameter Tuning

thread

1

```
Trial 0 starts-----
[I 2023-12-13 02:49:02,477] Trial 0 finished
Trial 1 starts-----
[I 2023-12-13 02:52:21,330] Trial 1 finished
Trial 2 starts-----
[I 2023-12-13 02:55:39,138] Trial 2 finished
Trial 3 starts-----
[I 2023-12-13 02:59:02,808] Trial 3 finished
Trial 4 starts-----
[I 2023-12-13 03:02:25,692] Trial 4 finished
Trial 5 starts-----
[I 2023-12-13 03:05:50,770] Trial 5 finished
Trial 6 starts-----
[I 2023-12-13 03:09:11,952] Trial 6 finished
Trial 7 starts-----
[I 2023-12-13 03:12:33,509] Trial 7 finished
Elapsed time: 1633.419298171997 s
```

threads

2

```
Trial 9 starts-----
Trial 8 starts-----
[I 2023-12-13 03:16:45,775] Trial 9 finished
[I 2023-12-13 03:16:45,776] Trial 8 finished
Trial 10 starts-----
Trial 11 starts-----
[I 2023-12-13 03:20:42,194] Trial 11 finished
[I 2023-12-13 03:20:42,197] Trial 10 finished
Trial 12 starts-----
Trial 13 starts-----
[I 2023-12-13 03:24:35,779] Trial 12 finished
[I 2023-12-13 03:24:35,780] Trial 13 finished
Trial 14 starts-----
Trial 15 starts-----
[I 2023-12-13 03:28:37,332] Trial 15 finished
[I 2023-12-13 03:28:37,334] Trial 14 finished
Elapsed time: 954.1966986656189 s
```

4 Result & Analysis

4.3 Hyperparameter Tuning

threads

4

```
Trial 17 starts-----
Trial 16 starts-----
Trial 18 starts-----
Trial 19 starts-----
[I 2023-12-13 03:33:46,613] Trial 16 finished
[I 2023-12-13 03:33:46,616] Trial 18 finished
[I 2023-12-13 03:33:46,617] Trial 19 finished
[I 2023-12-13 03:33:46,617] Trial 17 finished
Trial 20 starts-----
Trial 21 starts-----
Trial 22 starts-----
Trial 23 starts-----
[I 2023-12-13 03:38:36,570] Trial 23 finished
[I 2023-12-13 03:38:36,570] Trial 21 finished
[I 2023-12-13 03:38:36,571] Trial 22 finished
[I 2023-12-13 03:38:36,572] Trial 20 finished
Elapsed time: 591.6575014591217 s
```

threads

8


```
Trial 25 starts-----
Trial 26 starts-----
Trial 24 starts-----
Trial 27 starts-----
Trial 28 starts-----
Trial 29 starts-----
Trial 30 starts-----
Trial 31 starts-----
[I 2023-12-13 03:47:09,112] Trial 28 finished
[I 2023-12-13 03:47:09,113] Trial 31 finished
[I 2023-12-13 03:47:09,115] Trial 30 finished
[I 2023-12-13 03:47:09,116] Trial 29 finished
[I 2023-12-13 03:47:09,743] Trial 25 finished
[I 2023-12-13 03:47:12,015] Trial 27 finished
[I 2023-12-13 03:47:12,171] Trial 24 finished
[I 2023-12-13 03:47:12,826] Trial 26 finished
Elapsed time: 508.36329221725464 s
```

4 Result & Analysis

4.3 Hyperparameter Tuning

Count of Threads	Total Elapsed time (s)	speedup
Serial (1 thread)	1633.4193	1
2 threads	954.1967	1.7118
4 threads	591.6575	2.7608
8 threads	508.3633	3.2131



Multi-thread Tuning Elapsed Time and Speedup



4 Result & Analysis


4.4 Final Training & Model Evaluation

Training acc: 68.43%
Validation acc: 67.36%



4 Result & Analysis

4.4 Final Training & Model Evaluation



	precision	recall	f1-score	support
Apple__Apple_scab	0.64	0.58	0.61	504
Apple__Black_rot	0.63	0.60	0.62	497
Apple__Cedar_apple_rust	0.69	0.68	0.68	440
Apple__healthy	0.59	0.67	0.63	502
Blueberry__healthy	0.74	0.78	0.76	454
Cherry_(including_sour)__Powdery_mildew	0.70	0.76	0.73	421
Cherry_(including_sour)__healthy	0.74	0.77	0.76	456
Corn_(maize)__Cercospora_leaf_spotGray_leaf_spot	0.74	0.62	0.67	410
Corn_(maize)__Common_rust_	0.96	0.97	0.96	477
Corn_(maize)__Northern_Leaf_Blight	0.74	0.82	0.78	477
Corn_(maize)__healthy	0.94	0.95	0.95	465
Grape__Black_rot	0.52	0.56	0.54	472
Grape__Esca_(Black_Measles)	0.63	0.66	0.64	480
Grape__Leaf_blight_(Isariopsis_Leaf_Spot)	0.75	0.76	0.76	430
Grape__healthy	0.75	0.72	0.73	423
Orange__Haunglongbing_(Citrus_greening)	0.85	0.82	0.83	503
Peach__Bacterial_spot	0.69	0.61	0.64	459
Peach__healthy	0.72	0.85	0.78	432
Pepper,_bell__Bacterial_spot	0.51	0.51	0.51	478
Pepper,_bell__healthy	0.53	0.45	0.49	497
Potato__Early_blight	0.71	0.75	0.73	485
Potato__Late_blight	0.52	0.56	0.54	485
Potato__healthy	0.55	0.54	0.55	456
Raspberry__healthy	0.65	0.68	0.67	445
Soybean__healthy	0.74	0.75	0.74	505
Squash__Powdery_mildew	0.85	0.88	0.87	434
Strawberry__Leaf_scorch	0.67	0.68	0.68	444
Strawberry__healthy	0.71	0.72	0.72	456
Tomato__Bacterial_spot	0.69	0.75	0.72	425
Tomato__Early_blight	0.49	0.43	0.45	480
Tomato__Late_blight	0.49	0.38	0.43	463
Tomato__Leaf_Mold	0.59	0.55	0.57	470
Tomato__Septoria_leaf_spot	0.42	0.38	0.40	436
Tomato__Spider_mitesTwo-spotted_spider_mite	0.55	0.55	0.55	435
Tomato__Target_Spot	0.55	0.50	0.52	457
Tomato__Tomato_Yellow_Leaf_Curl_Virus	0.63	0.69	0.66	490
Tomato__Tomato_mosaic_virus	0.78	0.81	0.79	448
Tomato__healthy	0.87	0.87	0.87	481

Test acc: 67.33%

5 Conclusion

In data preprocessing

- Both frameworks effectively harness parallel processing to expedite computations and enhance overall efficiency.
- Compared to Dask, PyTorch exhibits a more substantial time reduction compared to Dask during data preprocessing tasks.

In model training

- Utilizing multiprocessing yields speedup improvements, but diminishing returns occur when processes exceed CPU core count.
- Increased GPU count with DataParallel accelerates training times, though multiprocessing outperforms in terms of performance.

In hyperparameters tuning

- Optuna demonstrates notable speedup improvements in hyperparameter tuning with an increasing number of threads.

References

- Gula, L. T. (2023, February 6). Researchers helping protect crops from pests. National Institute of Food and Agriculture. <https://www.nifa.usda.gov/about-nifa/blogs/researchers-helping-protect-crops-pests>
- Ingle, A. (2020, December 15). Plant disease classification - resnet-99.2%. Kaggle. <https://www.kaggle.com/code/atharvaingle/plant-disease-classification-resnet-99-2>
- Brownlee, J. (2019, July 5). How to manually scale image pixel data for deep learning. MachineLearningMastery. <https://machinelearningmastery.com/how-to-manually-scale-image-pixel-data-for-deep-learning/>
- Schmitt, M. (n.d.). Understanding dask architecture: Client, scheduler, workers. Datarevenue. <https://www.datarevenue.com/en-blog/understanding-dask-architecture-client-scheduler-workers>.
- Modi, R. (2022, February 25). ResNet - understand and implement from scratch. Medium. <https://medium.com/analytics-vidhya/resnet-understand-and-implement-from-scratch-d0eb9725e0db>
- Datasets & dataloaders. Datasets & DataLoaders - PyTorch Tutorials 2.1.1+cu121 documentation. (n.d.). https://pytorch.org/tutorials/beginner/basics/data_tutorial.html
- Zhu, J. (n.d.). Getting started with distributed data parallel. Getting Started with Distributed Data Parallel - PyTorch Tutorials 2.2.0+cu121 documentation. https://pytorch.org/tutorials/intermediate/ddp_tutorial.html

